RENDERING LEGIBLE THE HISTORY OF AMERICAN CAPITALISM

SEQUENTIAL HANDWRITTEN TEXT RECOGNITION WITH CRNN-CTC NETWORK

MOTIVATION

- In 2019, the New York Public Library digitized over 110,000 pages from the Brown Brothers & Company papers. A rich source of information on the workings of New York City finance, the history of American capitalism, and the role of "Northern" banks in the transatlantic slave trade, the Brown Brothers Collection has been vastly underutilized, in part because it is written in nineteenth-century longhand.

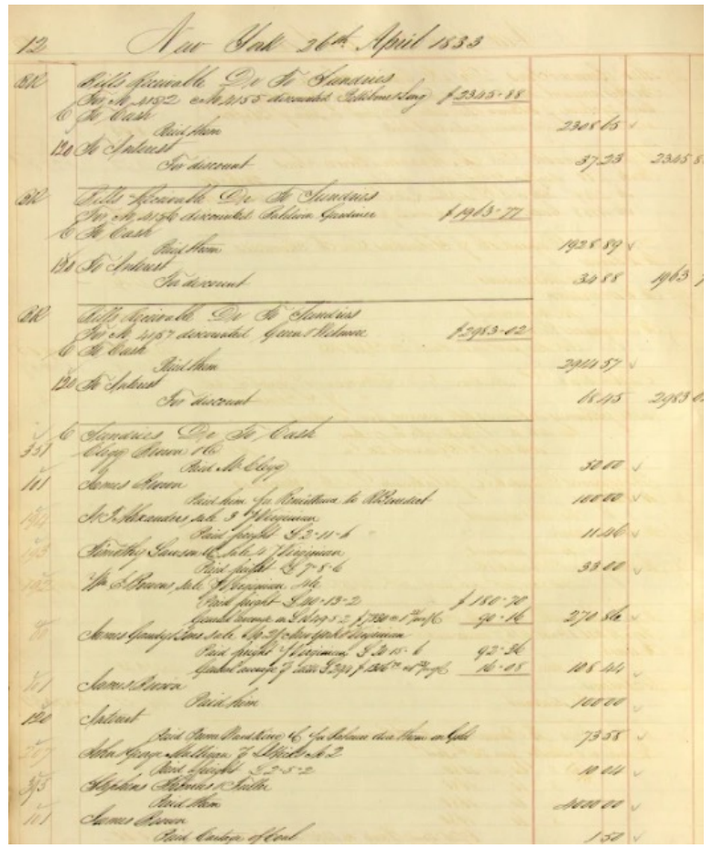

example ledger image

MACHINE LEARNING PIPELINE

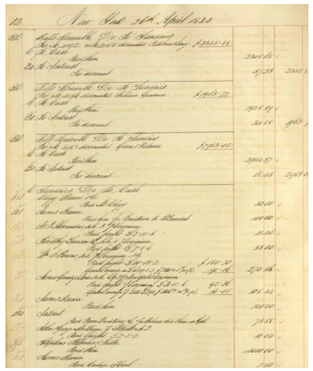

Input Data

Regression

Classification

Model

Supervised

Unsupervised

… the history of American capitalism, and the role of "Northern" banks in the transatlantic slave trade …

Prediction

DATASET

1. How to obtain the dataset?
2. How to convert the dataset into something that the machine can better understand?
3. What if the dataset is not large enough?

DATASET

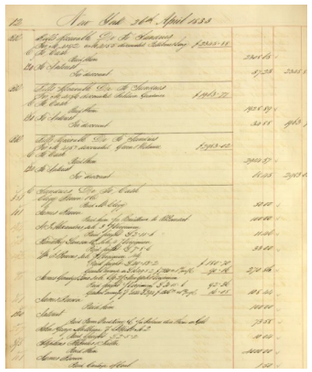

amazon

segmentation tool

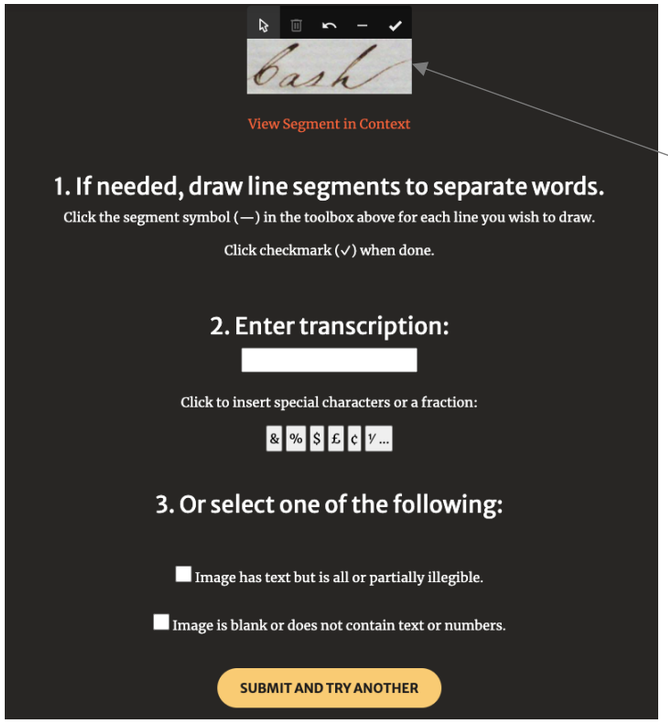

https://brownbros.newyorkscapes.org

credit to Grace Afsari-Mamagani

transcribers

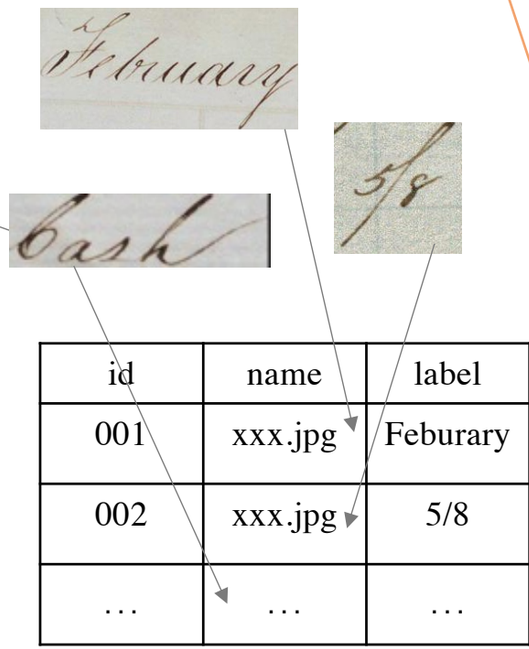

id	name	label
001	xxx.jpg	Feburary
002	xxx.jpg	5/8
...	...	...

DATA PREPROCESSING

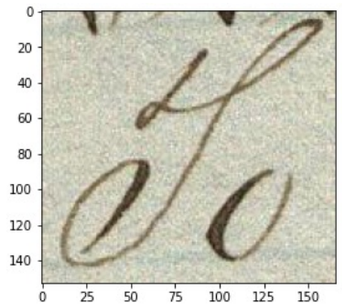

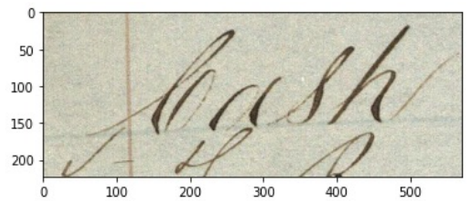

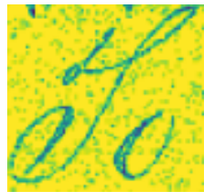

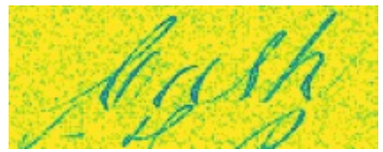

anti-clockwise rotation

Raw image

Thresholding + Normalization

Synthetic image

MODEL

1. What is the model?
2. What is the loss function?
3. How to evaluate the model?

VANILLA NERUAL NETWORK

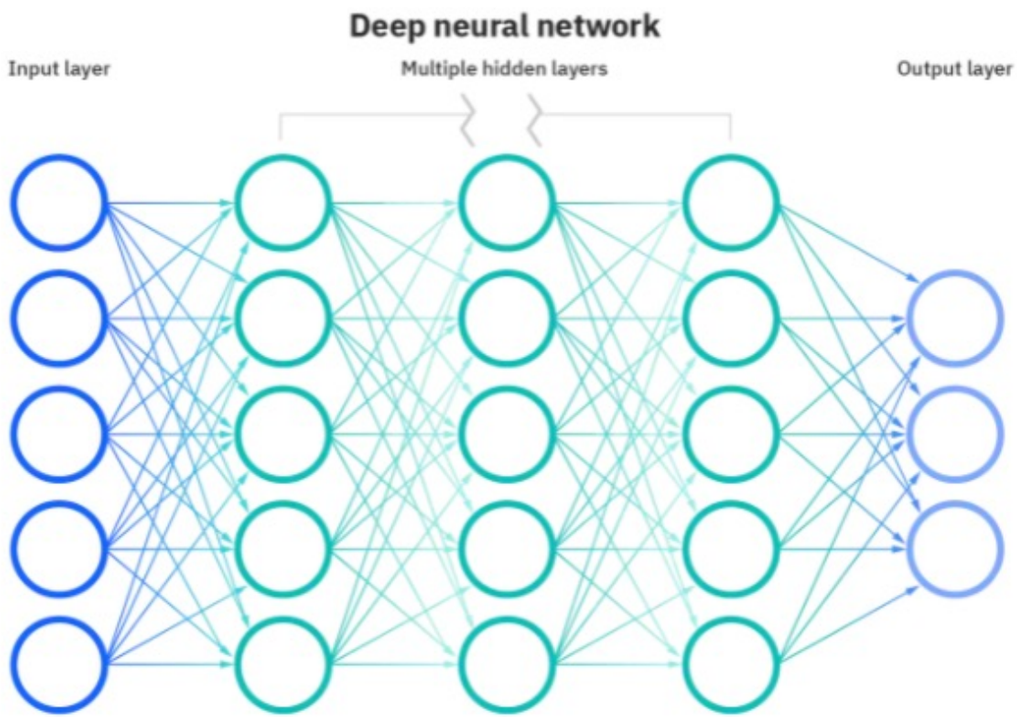

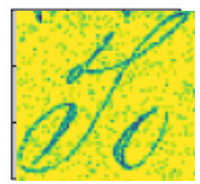

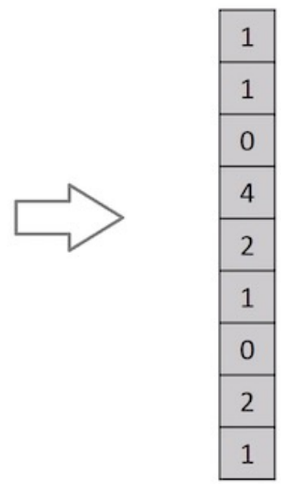

CONVOLUTIONAL NEURAL NETWORK

Capture the Spatial dependencies in an image

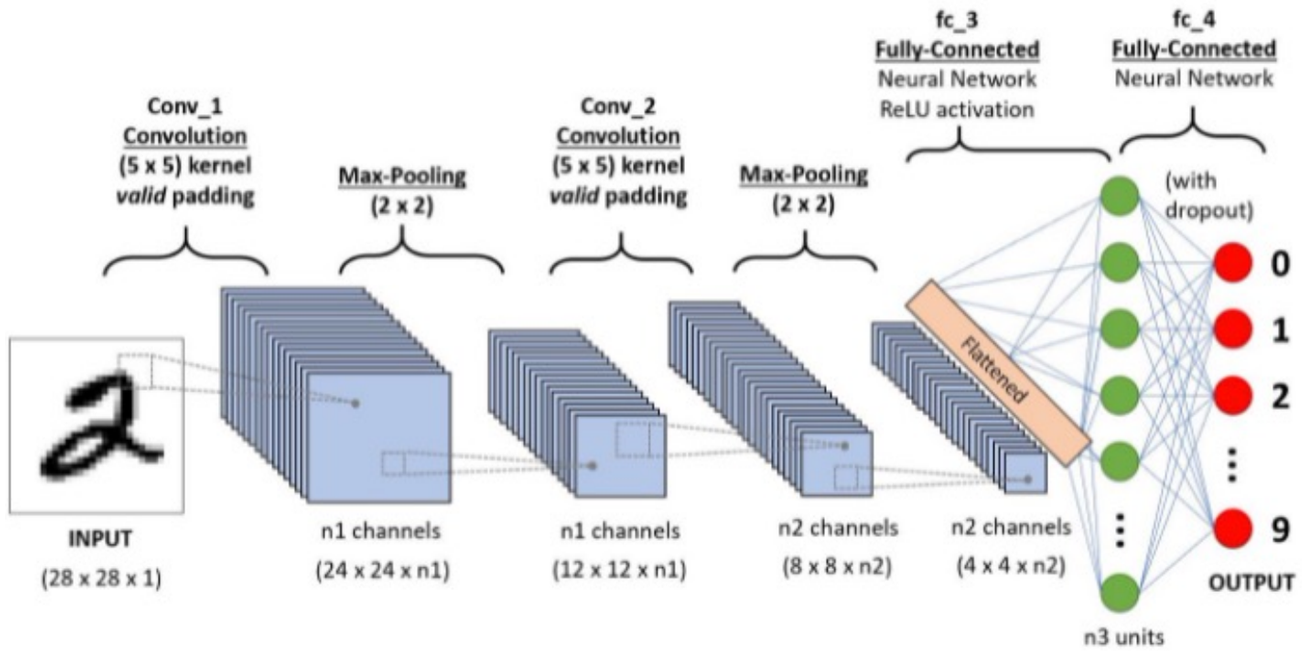

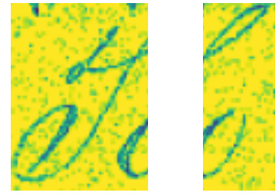

T *o*

RECURRENT NEURAL NETWORK

Capture the temporal dependencies in a sequence

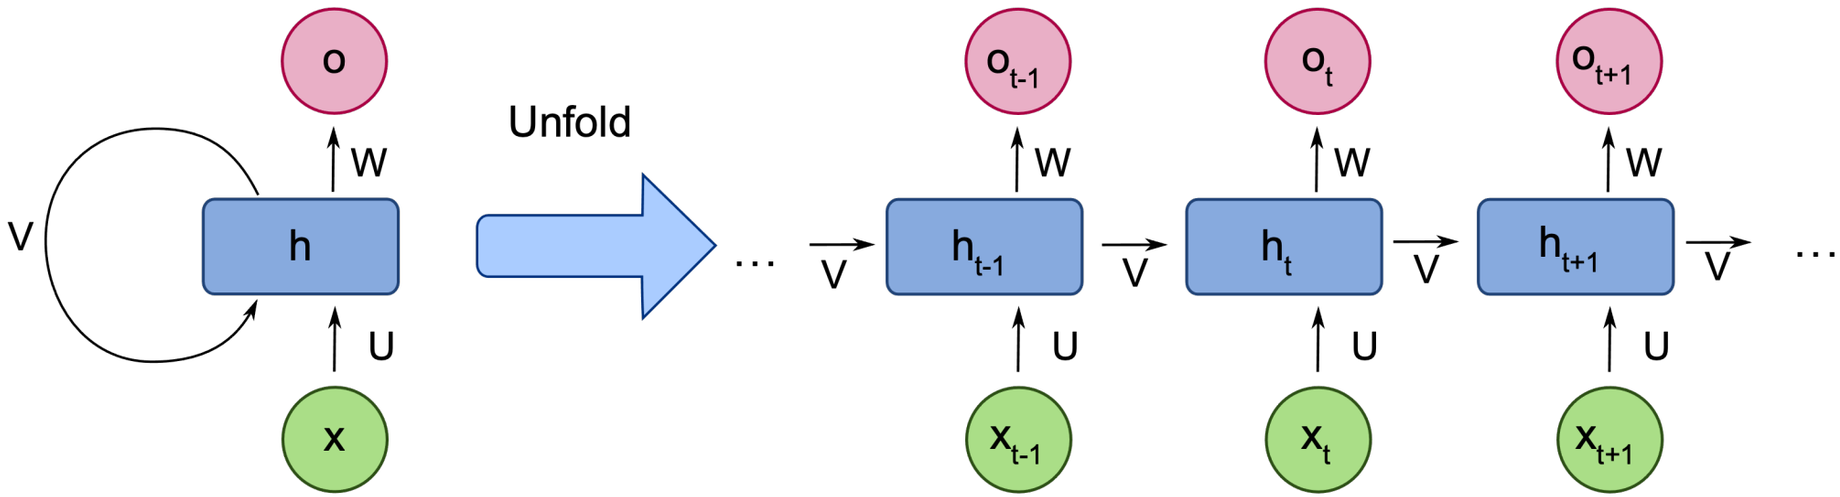

MODEL: CRNN-CTC

- Challenge: The output is a sequence of letters and the lengths may vary drastically.
- Use CNN to extract visual features + LSTM to capture sequential dependencies

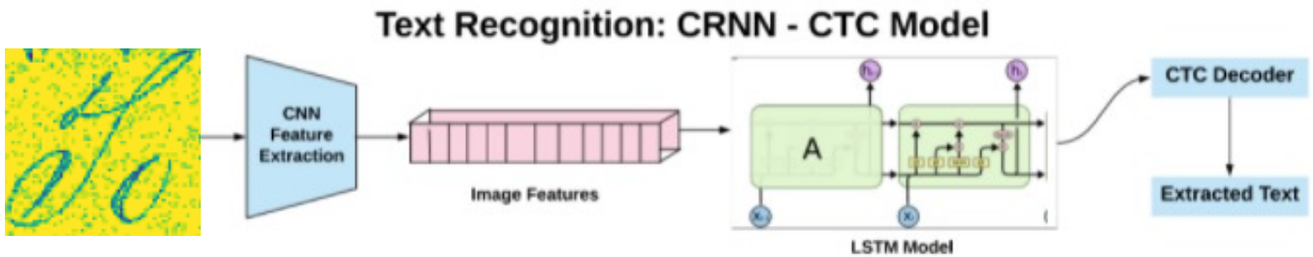

conv2 (Conv2D)	(None, 64, 256, 64)	18496
batch_normalization_1 (Batch	(None, 64, 256, 64)	256
activation_1 (Activation)	(None, 64, 256, 64)	0
max2 (MaxPooling2D)	(None, 64, 128, 64)	0
dropout_1 (Dropout)	(None, 64, 128, 64)	0

lstm1 (Bidirectional)	(None, 64, 512)	657408
lstm2 (Bidirectional)	(None, 64, 512)	1574912
lstm3 (Bidirectional)	(None, 64, 512)	1574912
lstm4 (Bidirectional)	(None, 64, 512)	1574912

figure adapted from Weights & Biases Articles

CONNECTIONIST TEMPORAL CLASSIFICATION (CTC)

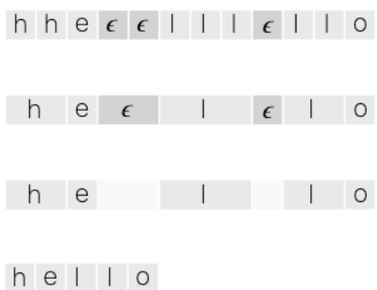

alignment

First, merge repeat characters.

Then, remove any ϵ tokens.

The remaining characters are the output.

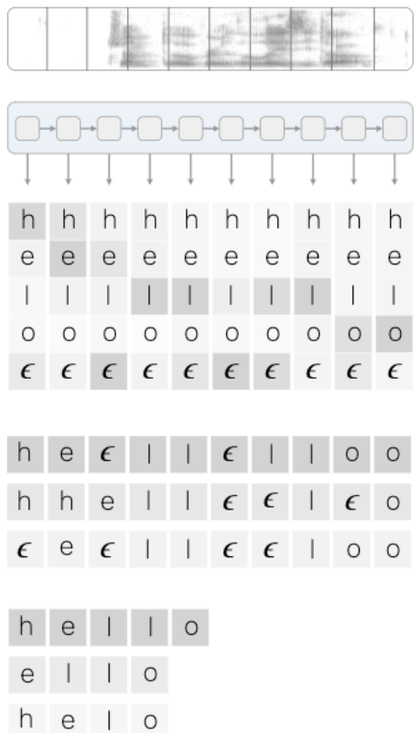

We start with an input sequence, like a spectrogram of audio.

The input is fed into an RNN, for example.

The network gives $p_t(a \mid X)$, a distribution over the outputs {h, e, l, o, ϵ} for each input step.

With the per time-step output distribution, we compute the probability of different sequences

By marginalizing over alignments, we get a distribution over outputs

loss

figure from Distill

METRICS

1. Word-level accuracy
2. Character-level accuracy
3. Edit distance
   - The Levenshtein distance allows deletion, insertion and substitution.
   - The Longest common subsequence (LCS) distance allows only insertion and deletion, not substitution.
   - The Hamming distance allows only substitution, hence, it only applies to strings of the same length.
   - The Damerau–Levenshtein distance allows insertion, deletion, substitution, and the transposition of two adjacent characters.
   - The Jaro distance allows only transposition.

1. **k**itten → **s**itten (substitute "s" for "k")
2. sitt**e**n → sitt**i**n (substitute "i" for "e")
3. sittin → sittin**g** (insert "g" at the end)

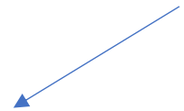

RESULTS

- Improve the word-level accuracy from 9% to 30.38%, char-level accuracy from 11% to 36.39%.

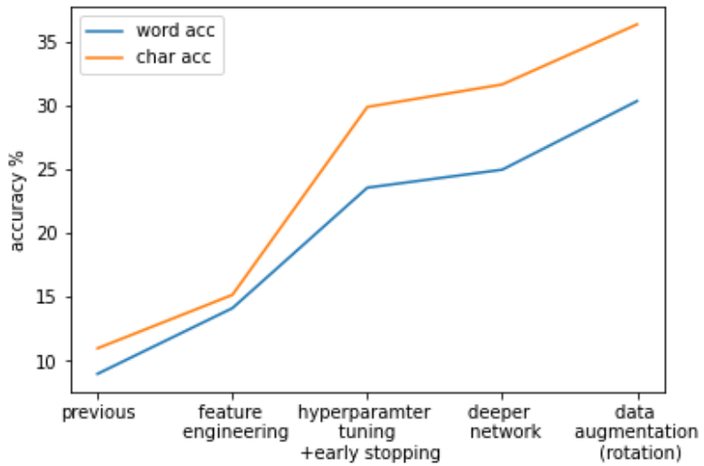

Correct examples

(short and/or frequent words)

to

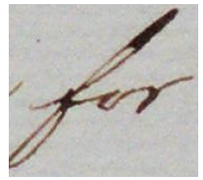

for

Sundries

Orleans

Incorrect examples

(Long, infrequent, blurry words)

Plantation (pred: lPlantation)

Price (pred: Brer)

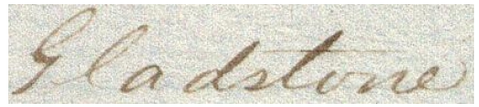

Gladstone (pred: Cawan)

CHALLENGE & FUTURE PLAN

1. The current preprocessed image still have a lot of noise
   - find a better preprocessing method to denoise, or
   - add some noise to the synthetic data to mimic the real-world images?
2. Numbers are important for analyzing financial activities, but the accuracy for them are low due to the lack of data (only 10% of the train set has numeric labels)
   - Only 10% of data has been transcribed now. Undergrad researchers are still working on transcribing more data this term!
   - Use synthetic images

- 3. Explore different models (e.g. Transformers)

RESOURCES

High Performance Computing

Additive Manufacturing and 3D Digitization

High Performance Computing

- High Performance Computing (NYU IT)
- HPC Research Project Space
- Research Cloud

Research Data and Tools

High Performance Computing provides supercomputer access and supporting software for researchers who need powerful processing resources. This includes the Greene supercomputer, one of the fastest HPC resources in higher education.

High Performance Computing (NYU IT)

HPC Research Project Space

Research Cloud

Google colab

Both have GPUs!